



department of computer science  
faculty of sciences

# Internet Programming

---

Multiprocessing in Unix

# Introduction

---

- Distributed programming is about **processes which communicate** over a network.
  - Obvious requirement: Excellent knowledge on managing them locally
  
- If there is ANYTHING you don't understand, **scream!!**
  - Else you'll be lost in later lessons!

# Today's Agenda

---

- ▣ Managing Unix Processes
- ▣ Inter-Process Communication
- ▣ Posix Threads
- ▣ Java Threads

---

# Managing Unix Processes

---

# Programs vs. Processes

---

- ❑ Computers can execute **programs**
- ❑ A **process** is one **instance** of a program, while it is executing
  - For example: I can run a simulation for my research
  - At the same time, there are many other programs executing: they are other processes
- ❑ The **same program** can be executed **multiple times** in parallel
  - Somebody else may log on my computer and start that simulator too
  - These are several **separate processes**, executing the **same program**
- ❑ A process can only be created by another process
  - E.g., when I type a command, my Unix shell process will create a new process to execute it

# Inside a Process

- ❑ One process is made of:
  - ❑ One executing **program** (i.e., a file containing the code to run)
  - ❑ **PID**: Process identifier (an integer)
  - ❑ **Memory** used to execute the program (text, data, heap, stack)
  - ❑ **PC**: Program counter
    - indicates where in the program the process currently is
  - ❑ A number of **signal handlers**
    - tells the program what to do when receiving signals
  
- ❑ One process can determine its own PID:

```
#include <sys/types.h>
#include <unistd.h>
pid_t getpid(void);
```

- ❑ Or the PID of its parent:

```
pid_t getppid(void);
```

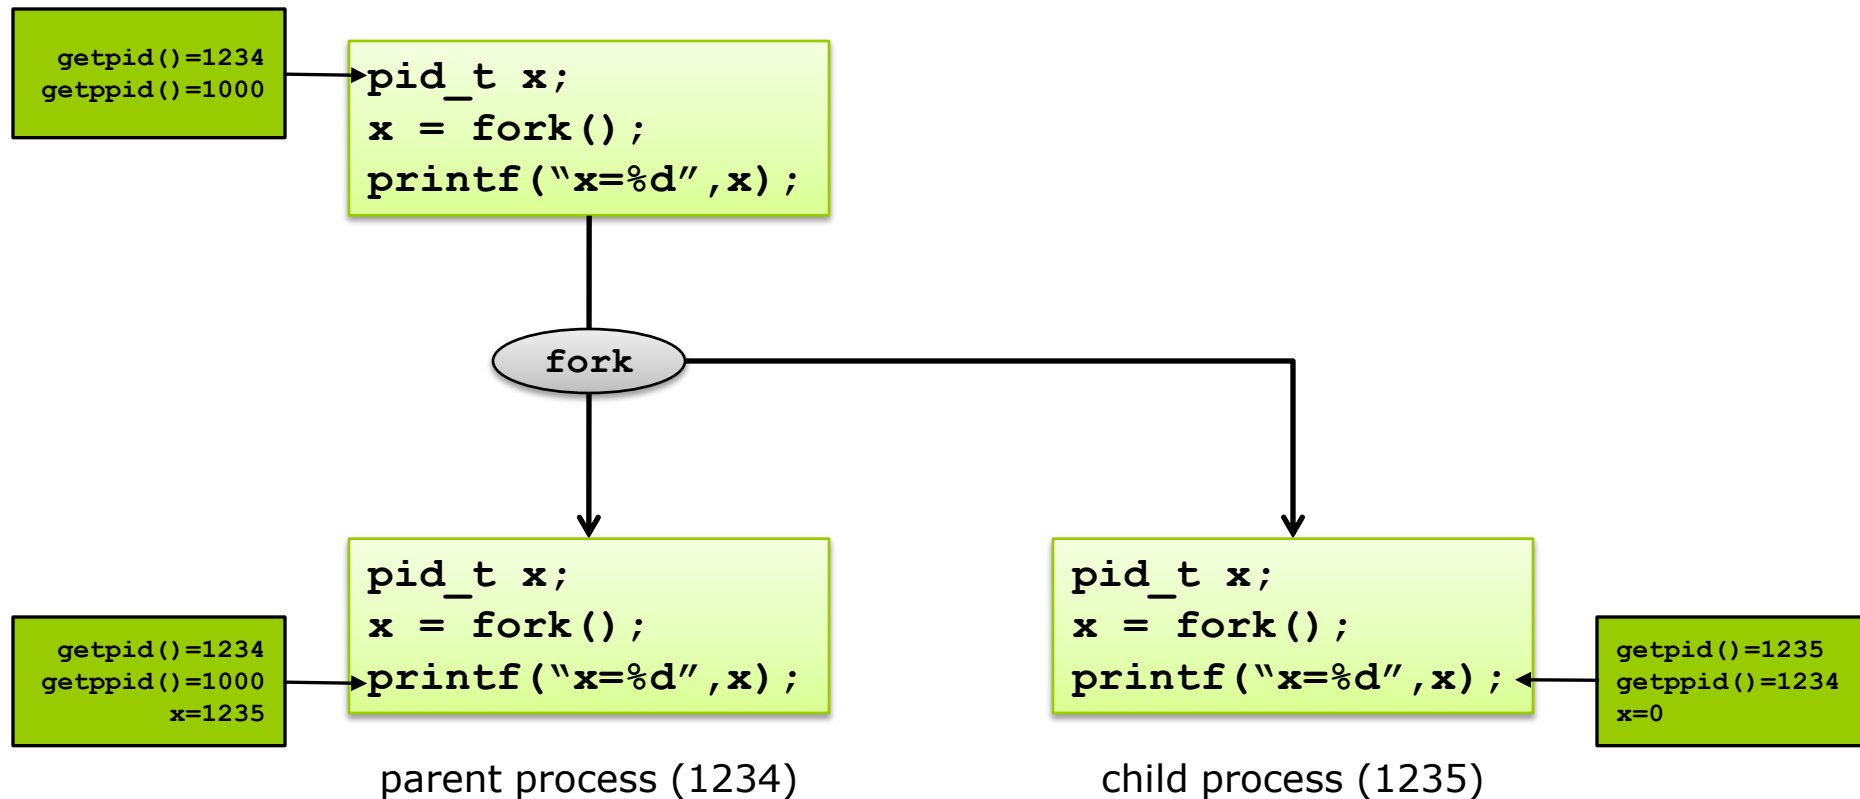
# The fork() System Call [1/3]

- In Unix systems, there is exactly **one way** to create a new process:

```
#include <sys/types.h>
#include <unistd.h>
pid_t fork(void);
```

- The child process is an **exact copy** of the parent
  - It is running the **same program**
  - Its **memory area** is an exact copy of the parent's memory
  - **Signal handlers** and **file descriptors** are copied too
  - Its **program counter (PC)** is at the same position within the program
    - I.e., just after the fork() call
- There is **one way** of distinguishing between the two processes:
  - **fork()** returns 0 to the child process
  - **fork()** returns the child's PID to the parent process
  - **fork()** returns -1 in case of error
- Most often, programs need to check the return value of **fork()**.

# The fork() System Call [2/3]





# The fork() System Call [3/3]

- Typically, processes diverge after fork() returns:

```
pid_t pid;
pid = fork();

if (pid<0) {perror("Fork error"); exit(1);}

if (pid==0) /* child */
{
    printf("I am the child process\n");
    while (1) putchar('c');
}
else /* parent */
{
    printf("I am the parent process\n");
    while(1) putchar('p');
}
```

# The Fork of Death!

---

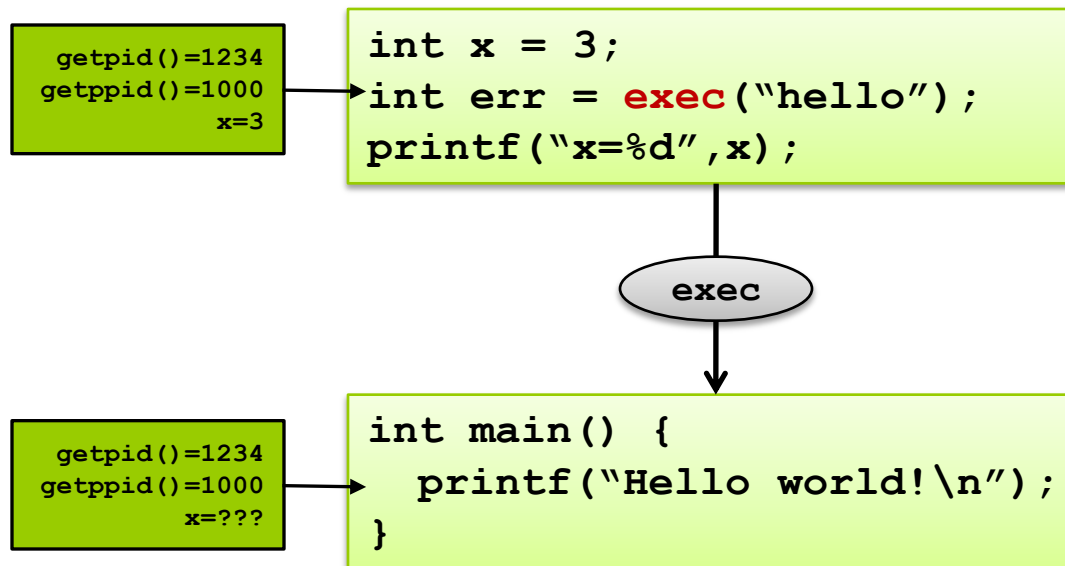
- ❑ You must be careful when using fork()!!
  - Very easy to create very harmful programs

```
while (1) {  
    fork();  
}
```

- ❑ Finite number of processes allowed
  - PIDs are 16-bit integers → maximum 65536 processes!
- ❑ Administrators typically take precautions
  - Limited quota on the number of processes per user
  - Make such experiments only at home 😊

# The exec() System Call [1/4]

- ❑ **exec()** allows a process to switch from one program to another
  - Code/Data for that process are destroyed
  - Environment variables are kept the same
  - File descriptors are kept the same
  - New program's code is loaded, then started (from the beginning)
  - There is no way to return to the previously executed program!



# The exec() System Call [2/4]

- fork() and exec() could be used separately
  - Imagine examples when this could happen?
- But most commonly they are used together:

```
if (fork()==0)      /* child */
{
    exec("hello");  /* load & execute new program */
    perror("Error calling exec()!\n");
    exit(1);
}
else /* parent */
{
    ...
}
```

# The exec() System Call [3/4]

- ❑ Command-line arguments are passed as an array of strings:
  - first argument should be `argv[0]` (i.e., the program name)
  - last argument should be `NULL` (to denote the end of the array)
  
- ❑ If you do not respect these constraints, **the program behavior is unspecified!**
  - Very strange behavior sometimes...
  
- ❑ E.g., to execute `"ls -l /home/spyros"`,
  - the program name is `"/bin/ls"`
  - and the arguments array is  
`{"ls", "-l", "/home/spyros", NULL}`

# The exec() System Call [4/4]

- There are 4 versions of exec(), depending on 2 criteria
  - Search for program in (i) absolute path, or (ii) \$PATH
  - Command-line arguments passed as list or single vector (array)

	absolute path	\$PATH
list ( <i>n+1</i> parameters)	execl()	execlp()
vector (2 parameters)	execv()	execvp()

- execl() example:

```
execl("/bin/ls", "ls", "-l", "/home/spyros", NULL);
```

- execv() example:

```
char *params[4] = {"ls", "-l", "/home/spyros", NULL};  
execv("/bin/ls", params);
```

# Questions

---

- ❑ Q1: What happens when you type a command in the shell?
  - The shell process **forks** (so, makes a copy of itself)
  - The child process **execs** the required program
  
- ❑ Q2: What if the command contains pipes?  
e.g.: `ls | grep ".txt" | wc`
  - The shell forks once per process
  - Creates **pipes** and arranges output of  $P_i$  to be input of  $P_{i+1}$
  
- ❑ Q3: How does a process end?
  - See next slides!

# Stopping a Process

---

- ❑ A process stops (without error) when:
  - The `main()` function returns
    - or -
  - the program calls `exit()`
  
- ❑ No process can directly kill another process, not even the kernel!
  - It can only send a **signal** to it
    - ❑ ...asking it to terminate itself
  - The signal will invoke the appropriate **signal handler** (a function)
  - By default, the signal **SIGINT** stops the process
    - ❑ This is what happens when you type ^C
    - ❑ The handler for that signal can be overwritten
  - Signal **SIGKILL** has the same effect
    - ❑ But it cannot be overwritten



# Signal Handling [1/3]

- Signals can be sent by one process (including a kernel process) to another
  - **SIGSEGV**: segmentation fault (non-authorized memory access)
  - **SIGBUS**: bus error (non-aligned memory access)
  - **SIGPIPE**: you tried to write in a pipe with no reader
  - **SIGCHLD**: one of your children processes has stopped
  - **SIGSTOP, SIGCONT**: pause and continue a process
  - **SIGUSR1, SIGUSR2**: two generic signals to be used by user programs
  - You can get a complete list of signals with **kill -l**

- Signals can also be sent explicitly: kill()

```
#include <sys/types.h>
#include <signal.h>
int kill(pid_t pid, int sig);
```

- **This function does not necessarily kill processes!**

# Signal Handling [2/3]

- To setup a custom signal handler:

```
#include <signal.h>
```

```
void (*signal(int signum, void (*sighandler)(int)))(int);
```

informally: `HANDLER *signal(int signum, HANDLER *sighandler);`

where HANDLER would be a function like this: `void sighandler(int)`

- Example:

```
#include <signal.h>

void myhandler(int sig) {
    printf("I received signal number %d\n", sig);
    signal(sig, myhandler); /* set it again it */
}

int main() {
    void (*oldhandler)(int);
    oldhandler = signal(SIGINT, myhandler);
    signal(SIGUSR1, myhandler);
    ...
}
```

# Signal Handling [3/3]

---

- ❑ When a signal handler is run, the signal gets associated back to its default behavior
  - You must call `signal()` **again** in the signal handler
  
- ❑ Attention: Unix normally does not queue signals
  - If the same signal is sent multiple times to the same process at short time intervals, **the signal may be delivered only once**
  
- ❑ Recent Unix systems have other signal handling calls that queue signals
  - We will not study them here

# Zombie Processes

---

- When a process terminates, it is not immediately removed from the system
  - The process which created it may be interested in its return value
  
- It gets to zombie state:
  - Its memory and resources are freed
  - It stays in the process table until its parent has received its termination status
  - If the parent has already finished, it is adopted by process 1 (init)
  
- Zombie processes must be dealt with!
  - Otherwise we eventually run out of PIDs
  - Prevention: The administrator should limit the number of processes a user can spawn
  
- To allow a zombie to die, its parent must explicitly “wait” for its children to finish
  - It can block until one of its children has died
  - Or it can setup a signal handler for the SIGCHLD signal to be asynchronously notified

# Waiting for Children Processes [1/2]

- To block until some child process completes:

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *status);
```

- **wait()** waits for **any child** to complete
  - Also handles an already completed (zombie) child → returns instantly
  - Returns the PID of the completed child
  - **status** indicates the process' return status

# Waiting for Children Processes [2/2]

- `waitpid()` gives you more control:

```
pid_t waitpid(pid_t pid, int *status, int option);
```

- `pid` specifies which child to wait for (-1 means any)
- `option=WNOHANG` makes the call return immediately, if no child has already completed (otherwise use `option=0`)

- Example: a `SIGCHLD` signal handler

```
void sig_chld(int sig) {  
    pid_t pid;  
    int stat;  
    while ( (pid=waitpid(-1, &stat, WNOHANG)) > 0 ) {  
        printf("Child %d exited with status %d\n", pid, stat);  
    }  
    signal(sig, sig_chld);  
}
```

- **Question:** Why is the while statement necessary?

---

# Inter-Process Communication

---

# Inter-Process Communication

---

- ❑ By default, processes cannot influence each other
  - Executed in isolation
  - Different address spaces
  
- ❑ There are 4 mechanisms for processes of the **same computer** to communicate
  - **Signals:** Send a signal to another process (SIGINT, SIGKILL, etc.)
  - **Pipes:** Communication channel to transfer data
  - **Shared memory:** the same memory area accessible to multiple processes
  - **Semaphores:** perform synchronization (e.g., to regulate access to shared memory)
  
- ❑ All these IPC methods work only between processes of the **same computer!**



# Pipes [1/3]


- A pipe is a unidirectional communication channel between processes
  - Write data on one end – Read it at the other end
  - Bidirectional communication? Use TWO pipes.

- Creating a pipe

```
#include <unistd.h>
int pipe(int fd[2]);
```

- The return parameters are:

- **fd[0]** is the file descriptor for reading
- **fd[1]** is the file descriptor for writing






# Pipes [2/3]

- Pipes are often used in combination with `fork()`

```
int main() {
    int pid, fd[2];
    char buf[64];

    if (pipe(fd)<0) exit(1);

    pid = fork();
    if (pid==0)      /* child */
    {
        close(fd[0]); /* close reader */
        write(fd[1],"hello, world!",14);
    }
    else {           /* parent */
    {
        close(fd[1]); /* close writer */
        if (read(fd[0],buf,64) > 0)
            printf("Received: %s\n", buf);
        waitpid(pid,NULL,0);
    }
    }
}
```



# Pipes [3/3]

---

- ❑ Pipes are used for example for shell commands like:  
`sort foo | uniq | wc`
- ❑ Pipes can only link processes which have a common ancestor
  - Because children processes inherit the file descriptors from their parent
- ❑ What happens when two processes with no common ancestor want to communicate?
  - Named pipes (also called FIFO)
  - A special file behaves like a pipe
  - Assuming both processes agree on the pipe's filename, they can communicate
    - ❑ `mkfifo()` for creating a named pipe
    - ❑ `open()`, `read()`, `write()`, `close()`

# Shared Memory [1/5]

---

- ❑ Shared memory allows multiple processes to have direct access to the same memory area
  - They can interact through the shared memory segment
  
- ❑ Shared memory segments must be **created**, then **attached** to a process to be usable. Then, you must **detach** and **destroy** them.
  
- ❑ When using shared memory, you must be very careful about **race conditions**, and solve them using semaphores.
  - If you don't know what a race condition is, then:
    - ❑ There will be an example later in this course
    - ❑ Shouldn't you attend the Operating Systems course? 😊

# Shared Memory [2/5]

- To **create** a shared memory segment:

```
#include <sys/ipc.h>
#include <sys/shm.h>
int shmget(key_t key, int size, int shmflg);
```

- **key**: rendezvous point (key=IPC\_PRIVATE if it will be used by children processes)
- **size**: size of the segment in bytes
- **shmflg**: options (access control mask)
- **Return value**: a shm identifier (or -1 for error)

# Shared Memory [3/5]

- To **attach** a shared memory segment:

```
#include <sys/types.h>
#include <sys/shm.h>
void *shmat(int shmid, const void *shmaddr, int shmflg);
```

- **shmid**: shared memory identifier (returned by shmget)
- **shmaddr**: address where to attach the segment, or NULL if you don't care
- **shmflg**: options (access control mask)
  - SHM\_RDONLY: read-only

- To **detach** a shared memory segment:

```
int shmdt(const void *shmaddr);
```

- shmaddr: segment address
- **Attention**: shmdt() does **not** destroy the segment!

# Shared Memory [4/5]

- To **destroy** a shared memory segment:

```
#include <sys/ipc.h>
#include <sys/shm.h>
int shmctl(int shmid, int cmd, struct shm_id *buf);
```

- **shmid**: shared memory identifier (returned by shmget)
  - **cmd=IPC\_RMID** to destroy the segment
  - **buf**: NULL as far as we are concerned
- 
- Shared memory segments stay persistent even after all processes have died!
    - You **must** destroy them in your programs
    - The **ipcs** command shows existing segments (and semaphores)
    - You can destroy them by hand with: **ipcrm shm <id>**

# Shared Memory [5/5]

## □ Example:

```
int main() {
    int shmid = shmget(IPC_PRIVATE, sizeof(int), 0600);
    int *shared_int = (int *) shmat(shmid, 0, 0);
    *shared_int = 42;

    if (fork()==0) {
        printf("The value is: %d\n", *shared_int);
        *shared_int = 12;
        shmdt((void *) shared_int);
    }
    else {
        sleep(1);
        printf("The value is: %d\n", *shared_int);
        shmdt((void *) shared_int);
        shmctl(shmid, IPC_RMID, 0);
    }
}
```



# Race Conditions [1/2]

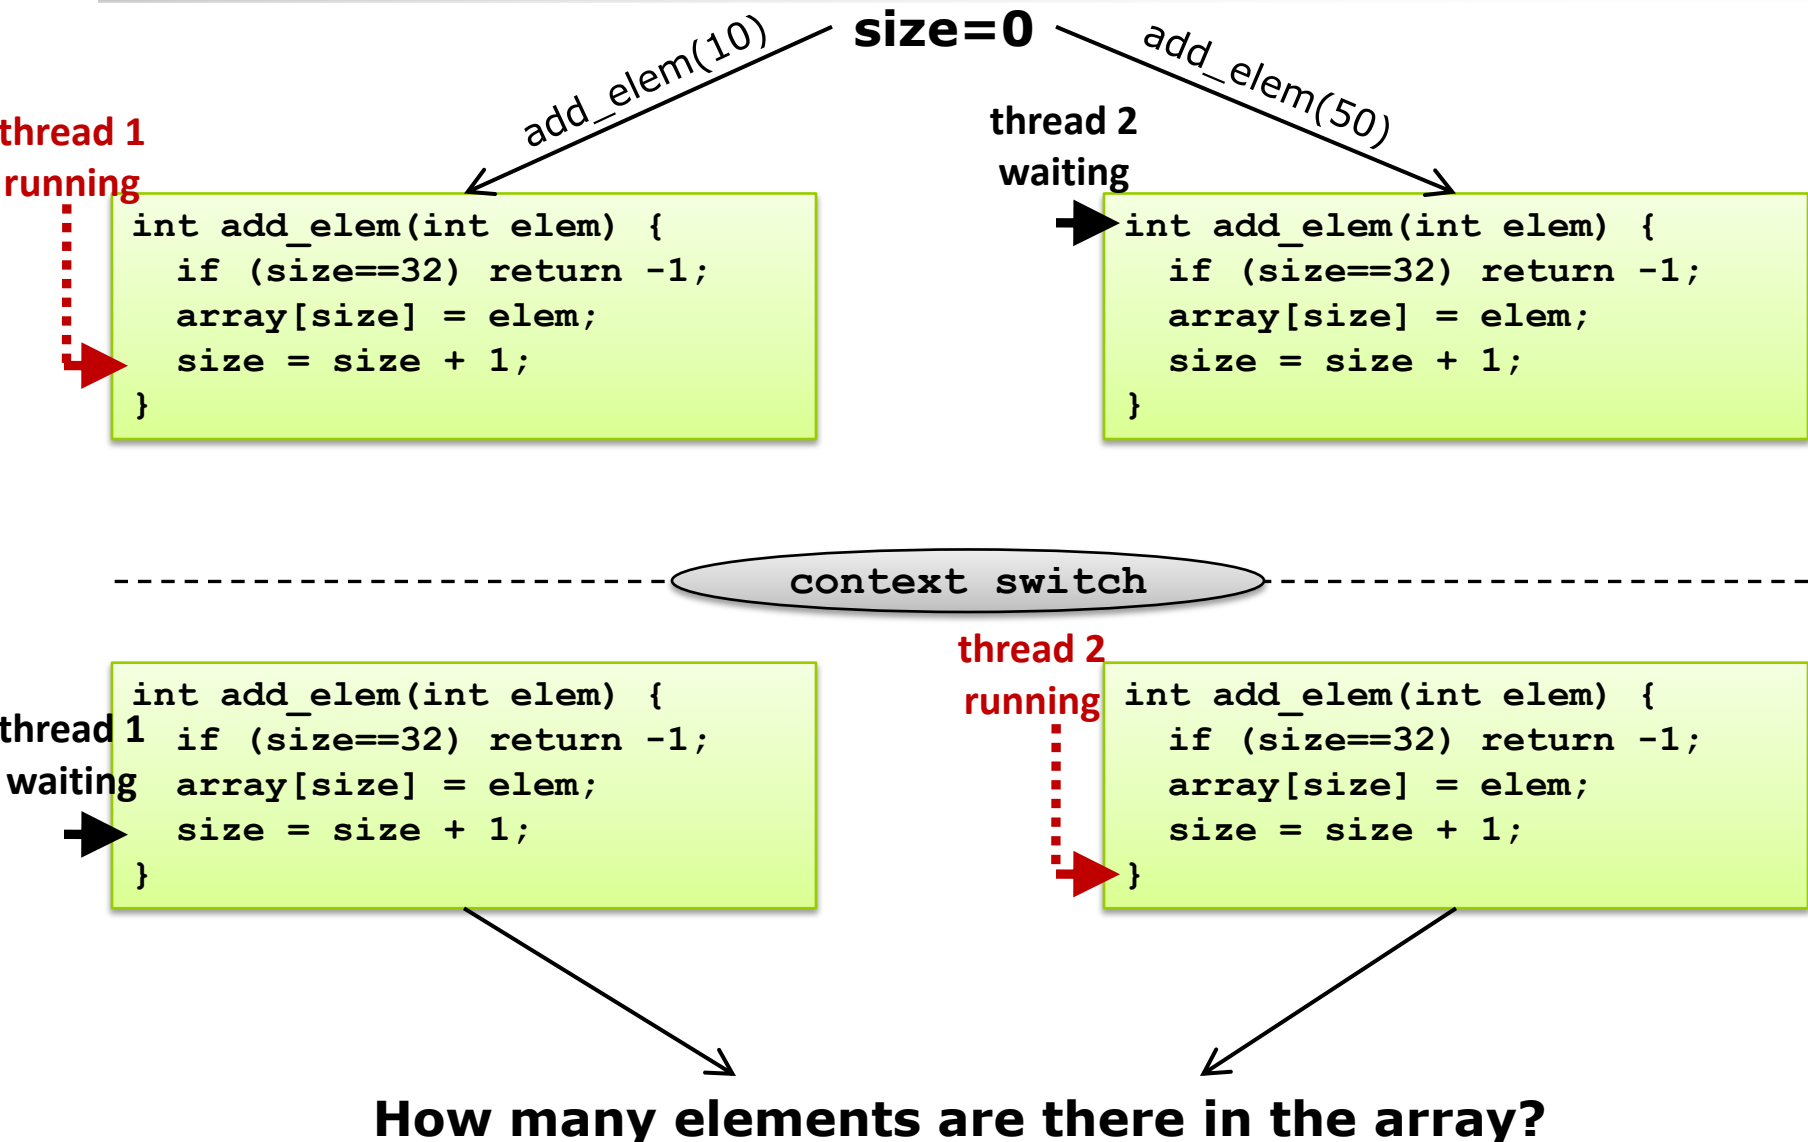
- When several processes share a resource, you **must always** wonder if synchronization is needed
- Example: Add an element to the end of an array

```
int array[32], size;

int add_elem(int elem)
{
    if (size==32) return -1;
    array[size] = elem;
    size = size + 1;
}
```

- This is a **race condition**: if two processes execute the function simultaneously, the result can be **wrong**

# Race Conditions [2/2]



# Semaphores [1/6]

---

- In many cases you need to synchronize processes
  - When they share a resource (shared memory, file descriptor, device, etc.)
  - When a process needs to wait for a given event
  
- Semaphores are positive integers with two methods: **UP()** and **DOWN()**
  - **DOWN():**
    - If  $\text{sem} \geq 1$  then  $\text{sem} = \text{sem} - 1$
    - Otherwise, block the process
  - **UP():**
    - If there are blocked processes, then unblock one of them
    - Otherwise,  $\text{sem} = \text{sem} + 1$
  - Semaphore operations are atomic

# Semaphores [2/6]

---

- Semaphores are generally used for two goals:
  - **Mutual exclusion:** only one process at a time can be within a section of code
    - Start with **sem=1**
    - To **enter** the mutex section: DOWN(sem)
    - To **leave** the mutex section: UP(sem)
  - **Process synchronization:** wait for a given event
    - Start with **sem=0**
    - To **wait** for the event: DOWN(sem)
    - To **trigger** the event: UP(sem)
    - What happens if the event is triggered before the other process waits for it?
- QUESTION: How can you solve the race condition in add\_elem()?
- QUESTION: What if a critical region should accept up to  $K > 1$  processes?

# Semaphores [3/6]

- Semaphores are created in arrays
- To **create** an array of semaphores:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
int semget(key_t key, int nsems, int semflg);
```

- **key**: rendezvous point or IPC\_PRIVATE
- **nsems**: number of semaphores to create
- **semflg**: access rights
- Return value: a semaphore array identifier

# Semaphores [4/6]

- UP() and DOWN() are realized with the same function:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
int semop(int semid, struct sembuf *sops, unsigned nsops);
```

- **semid**: the semaphore array identifier
- **sops**: an array of commands to be issued
- **nsops**: the size of sops

- Here's the sembuf struct:

```
struct sembuf {
    short sem_num; /* semaphore number: 0 = first */
    short sem_op; /* semaphore operation */
    short sem_flg; /* operation flags */
};
```

- **sem\_op**: 1 for UP(), -1 for DOWN()
- **All operations are executed together atomically as a whole**

# Semaphores [5/6]

## □ To manipulate a semaphore:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
int semctl(int semid, int semnum, int cmd, ...);
```

- **semid**: the semaphore array identifier
- **semnum**: the semaphore number
- **cmd**: command
  - IPC\_RMID: destroy the semaphore
  - GETVAL: return the value of the semaphore

- Like shared memory segments, semaphores persist in the system after all processes have completed
  - You must destroy them in your programs
  - By hand: `ipcs`, `ipcrm`

# Semaphores [6/6]

- A (rather meaningless) example:

```
int main() {
    struct sembuf up    = {0,  1, 0};
    struct sembuf down  = {0, -1, 0};
    int my_sem, v1, v2, v3, v4;

    my_sem = semget(IPC_PRIVATE, 1, 0600); /* create semaphore */
    v1 = semctl(my_sem, 0, GETVAL);

    semop(my_sem, &up, 1);                /* UP() */
    v2 = semctl(my_sem, 0, GETVAL);

    semop(my_sem, &down, 1);               /* DOWN() */
    v3 = semctl(my_sem, 0, GETVAL);

    semctl(my_sem, 0, IPC_RMID);           /* destroy */
    v4 = semctl(my_sem, 0, GETVAL);

    printf("Semaphore values: %d %d %d %d\n", v1, v2, v3, v4);
}
```



---

# Posix Threads

---

# Threads vs. Processes

---

- ❑ Multi-process programs are expensive:
  - `fork()` needs to copy all the process' memory, etc.
  - interprocess communication is hard
  
- ❑ Threads: “lightweight processes”
  - One process contains several “threads of execution”
  - All threads execute the same program (but can be at different stages within it)
  - All threads share process instructions, global memory, open files, and signal handlers
  - Each thread has its own thread ID, program counter (PC), stack and stack pointer (SP), `errno`, and signal mask
  - There are special synchronization primitives between threads of the same process

# Threads in C and Java

---

## □ Posix Threads

- Posix Threads (pthreads) are standard among Unix systems
  - Also available on Windows through 3<sup>rd</sup> party libraries (Pthreads-w32)
- The operating system must have special support for threads
  - Linux, Solaris, and virtually all Unix systems have it
- Programs must be linked with `-lpthread`
  - Beware: Solaris will compile fine even if you forget the `-lpthread` (but your program will not work)

## □ Java Threads

- Threads are a native feature of Java: every virtual machine has thread support
- They are portable on any Java platform
- Java threads can be:
  - mapped to operating system threads (kernel threads or native threads)
  - or emulated in user space (user threads or green threads)

# Creating a pthread

## □ To create a pthread:

```
#include <pthread.h>
int pthread_create(pthread_t *thread,
                  pthread_attr_t *attr,
                  void *(*start_routine)(void *),
                  void *arg);
```

- **thread**: thread id (this is a return argument)
- **attr**: attributes (i.e., options)
- **start\_routine**: function that the thread will execute
- **arg**: parameter to be passed to the thread
- Return value: 0 on success, error value on failure

## □ To initialize and destroy the default pthread attributes

```
int pthread_attr_init(pthread_attr_t *attr);
int pthread_attr_destroy(pthread_attr_t *attr);
```

# Stopping a pthread

- A pthread stops when:
  - Its process stops
  - Its parent thread stops
  - Its `start_routine()` function returns
  - It calls `pthread_exit`:

```
#include <pthread.h>
void pthread_exit(void *retval);
```

- Like processes, stopped threads must be waited for:

```
#include <pthread.h>
int pthread_join(pthread_t th, void **thread_return);
```

# pthread Create/Delete Example

- To create a pthread:

```
#include <pthread.h>

void *func(void *param) {
    int *p = (int *) param;
    printf("New thread: param=%d\n", *p);
    return NULL;
}

int main() {
    pthread_t id;
    pthread_attr_t attr;
    int x = 42;

    pthread_attr_init(&attr);
    pthread_create(&id, &attr, func, (void *) &x);
    pthread_join(id, NULL);
}
```

# Detached Threads

- A “detached” thread:
  - does not need to be joined by `pthread_join()`
  - does not stop when its parent thread stops
- By default, threads are “joinable” (i.e. “attached”)
- To create a detached thread, set an attribute before creating the thread:

```
pthread_t id;  
pthread_attr_t attr;  
  
pthread_attr_init(&attr);  
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);  
pthread_create(&id, &attr, func, NULL);
```

- You can also detach a thread later with `pthread_detach()`
  - But you cannot reattach it!

# Race Conditions with Threads

---

- ❑ Threads share most resources by default
  - memory , file descriptors, etc.
  
- ❑ **Attention! Danger!**
  - It is very easy to make race conditions without even noticing
  
- ❑ You must **always** wonder if synchronization is needed
  - And solve them with special thread synchronization primitives
  
- ❑ Pthreads have two synchronization concepts:
  - **Mutex**
  - **Condition Variables**



# Pthread Sync with Mutex [1/2]

## □ Mutex (mutual exclusion)

```
#include <pthread.h>
int pthread_mutexattr_init(pthread_mutexattr_t *attr);
int pthread_mutex_init(pthread_mutex_t *mutex,
                       const pthread_mutexattr_t *mutexattr);
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

# Pthread Sync with Mutex [2/2]

## □ Example

```
pthread_mutex_t mutex;

int add_elem(int elem) {
    int n;
    pthread_mutex_lock(&mutex);
    if (size==32) {
        pthread_mutex_unlock(&mutex);
        return -1;
    }
    array[size++] = elem;
    n = size;
    pthread_mutex_unlock(&mutex);
}

int main() {
    pthread_mutexattr_t attr;
    pthread_mutexattr_init(&attr);
    pthread_mutex_init(&mutex, &attr);
    ...
    pthread_mutex_destroy(&mutex);
}
```

# Thread Safety with Unix Primitives

---

- ❑ Attention: Some Unix primitives and library functions have similar internal race conditions
  - Many primitives were not designed with threads in mind
  - When writing a multi-threaded program, you **must always check** in the man pages **if library functions are thread-safe**
  
- ❑ Example:
  - `gethostbyname()` is not thread-safe
  - You can still use it in multi-threaded programs
    - ❑ **but only within a mutex protection**

# Pthread Sync with Condition Variables

---

- Condition Variables
  - One thread waits for an event triggered by another thread
  
- Condition variables always work together with a mutex and a predicate variable (e.g., a boolean, an int, etc.)
  - The predicate variable is used to make sure the event has really been triggered
    - Sometimes threads are woken up without the event having been triggered!
  - The mutex is used to synchronize access to the condition variable and the predicate variable
  
- Details are out of the scope of Internet Programming
  - You can find a detailed explanation at:
    - <http://www.ibm.com/developerworks/linux/library/l-posix3/>
    - <https://computing.llnl.gov/tutorials/pthreads/>

# API for Condition Variables

## □ At init time:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;  
int predicate = FALSE;  
  
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```

## □ To wait for an event:

```
pthread_mutex_lock(&mutex);  
while (predicate==FALSE)  
    pthread_cond_wait(&cond_var, &mutex);  
predicate = TRUE;  
pthread_mutex_unlock(&mutex);
```

## □ To trigger an event:

```
pthread_mutex_lock(&mutex);  
predicate = TRUE;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

---

# Java Threads

---

# Creating Java Threads

- ❑ A thread is a class that inherits from Thread
- ❑ You must overload its run() method:

```
public class MyThread extends Thread {  
    private int argument;  
  
    MyThread(int arg) {  
        argument = arg;  
    }  
  
    public void run() {  
        System.out.println("New thread started! arg=" + argument);  
    }  
}
```

- ❑ To start the thread:

```
MyThread t = new MyThread(1050);  
t.start();
```

# Stopping Java Threads

---

- ❑ A thread stops when its run() method returns
- ❑ You do not need to join() for a Java thread to finish
  - But you can, if you want:

```
MyThread t = new MyThread(1050);  
t.start();  
...  
t.join();
```



# Synchronization with Monitors [1/2]

- A monitor is similar to a mutex:

```
public class AnotherClass {  
    synchronized public void methodOne() { ... }  
    synchronized public void methodTwo() { ... }  
    public void methodThree()           { ... }  
}
```

- Each object contains one mutex, which is locked when entering a synchronized method and unlocked when leaving
- Locking is at the object (instance) level
  - Two threads cannot be executing synchronized methods of a given object at the same time
    - But this is allowed by the same thread (e.g., one synchronized method calls another)
  - No restrictions apply to different objects of a given class

# Synchronization with Monitors [2/2]

- So, the previous class:

```
public class AnotherClass {  
    synchronized public void methodOne() { ... }  
    synchronized public void methodTwo() { ... }  
    public void methodThree() { ... }  
}
```

- ...is equivalent to:

```
public class AnotherClass {  
    private Mutex mutex;  
    public void methodOne() { mutex.lock(); ...; mutex.unlock(); }  
    public void methodTwo() { mutex.lock(); ...; mutex.unlock(); }  
    public void methodThree() { ... }  
}
```

...except that the Mutex class does not exist!

- **QUESTION:** Can you implement a Mutex class using monitors?

# Condition Variables in Java [1/2]

- ❑ There is no real condition variable in Java
  - But you can explicitly block a thread
- ❑ All Java classes inherit from class Object the following:

```
class Object {  
    void wait();           /* blocks the calling thread */  
    void notify();        /* unblocks ONE thread blocked by this object */  
    void notifyAll();     /* unblocks ALL threads blocked by this object */  
}
```

- **wait()**: causes current thread to wait until another thread invokes the `notify()` or `notifyAll()` method for this object
- **notify()**: wakes up a single thread that is waiting on this object's monitor
- **notifyAll()**: wakes up all threads that are waiting on this object's monitor

# Condition Variables in Java [2/2]

- ❑ This means, each object contains exactly one (and no more) condition variable
- ❑ The `wait()`, `notify()` and `notifyAll()` methods can only be called inside a monitor of that same instance. E.g.:

```
Integer myObj; // or any other type of object

synchronized(myObj)
{
    myObj.wait();
}
```

## ❑ Questions:

- [1/3] What if you need several condition variables in a given object?
- [2/3] Can you implement a `ConditionVariable` class in Java?
- [3/3] Can you implement a `Mutex` class in Java?